



SAE Bottleneck Geometry – Full-Precision vs Quantized [d_sae = 512]

Per-Tensor


Full-Precision



8-bit



6-bit



4-bit



2-bit



Per-Feature

